

Strategic Blueprint for a Decentralized, Vision-Language Model-Driven Autonomous Driving Data Collection Platform in India

Executive Overview

The transition toward fully autonomous mobility and Advanced Driver Assistance Systems (ADAS) in emerging markets is currently obstructed by profound infrastructural and environmental friction. Standardized ADAS architectures, which are predominantly trained on highly structured, rule-compliant traffic scenarios prevalent in Western geographies, consistently fail when confronted with the heterogeneous, unpredictable, and chaotic nature of Indian traffic. The fundamental bottleneck preventing the localization of autonomous driving technology is not a deficit in computational power, but rather a critical absence of edge-case-rich, geographically specific training data.

The ensuing analysis delineates an exhaustive, multi-disciplinary strategy for developing a low-cost, edge-computing hardware device designed for mass deployment in consumer vehicles. This system is engineered to passively collect, filter, and transmit high-value driving data to train next-generation Vision-Language Models (VLMs) and end-to-end autonomous driving architectures. By transforming this proprietary dataset into a Business-to-Business (B2B) Software-as-a-Service (SaaS) offering, the proposed platform directly addresses the immediate data localization needs of automotive Original Equipment Manufacturers (OEMs). Crucially, the analysis thoroughly investigates the incentive structures required to drive widespread consumer adoption, proposing a triad of Decentralized Physical Infrastructure Network (DePIN) tokenomics, Usage-Based Insurance (UBI) integrations, and on-device safety utilities to compel daily commuters to voluntarily participate in this data-harvesting ecosystem. Furthermore, this report acts as an exhaustive architectural and commercial roadmap meticulously tailored for presentation at the Vihaan 9.0 IEEE Delhi Technological University (DTU) hackathon.

The Indian Autonomous Driving Data Deficit

The operational design domain of Indian roads presents a unique set of challenges that fundamentally break traditional perception and planning algorithms. Western autonomous vehicle companies rely heavily on deterministic models, high-definition 3D mapping, and highly structured, predictable environments.¹ Conversely, the Indian driving ecosystem is characterized by unstructured traffic flows, faded or entirely absent lane markings, arbitrary vehicular movements, and an incredibly diverse mix of traffic participants.² The intermingling of

handcarts, autorickshaws, stray livestock, and pedestrians with high-speed motorized traffic creates an environment that is less a predictable grid and more a fluid, probabilistic ecosystem.³

Limitations of Current ADAS Implementations

Consequently, camera-centric ADAS systems currently deployed in premium Indian vehicles often exhibit erratic behavior. Vehicles equipped with ADAS features, such as the Mahindra XUV700, frequently trigger false positives, abrupt emergency braking, or nervous steering corrections.⁶ These failures are largely driven by visual clutter, sudden and unpredictable maneuvers by two-wheelers, and harsh environmental conditions, including heavy dust and extreme glare, which blind standard optical sensors.³ Standard features like Lane Keep Assist (LKA) and Adaptive Cruise Control (ACC) struggle to maintain functionality when lane markings disappear or when chaotic traffic prevents the maintenance of a standard following distance.¹⁰

Industry leaders and researchers suggest that autonomous mobility in India will not begin with Level 5 robotaxis deployed seamlessly into urban centers. Instead, it will emerge through the gradual adaptation of systems capable of surviving the chaos of unstructured roads, starting in controlled environments or specific highway corridors.¹¹ The prevailing consensus dictates that if an autonomous system can successfully navigate the unpredictable edge cases of Indian traffic, the underlying technology possesses a universal robustness applicable anywhere in the world.¹

The Shift Toward Vision-Language Models (VLMs)

To achieve this level of robust autonomy, the industry is aggressively transitioning away from traditional Convolutional Neural Networks (CNNs) and rigid modular pipelines, moving toward End-to-End (E2E) autonomous driving powered by Vision-Language Models (VLMs). Existing E2E models are typically optimized to mimic driving patterns observed in data, but they lack the ability to capture underlying reasoning processes, which limits their effectiveness in novel edge cases.¹² VLMs, however, leverage large-scale image-text pretraining to demonstrate robust instruction-following and profound contextual reasoning capabilities.¹³

By utilizing VLMs, an autonomous system can process a complex, chaotic scene and output a reasoned driving decision in natural language before executing the control command.¹⁵ For example, a VLM can interpret a scenario involving a cow crossing a chaotic intersection while a traffic police officer utilizes localized hand gestures. The model utilizes a Chain-of-Thought (CoT) process to deduce the correct course of action, bridging the gap between raw visual perception and nuanced human-like reasoning.¹⁷ Training these sophisticated models, however, requires immense volumes of localized, multi-modal data encompassing diverse edge cases. The lack of such data creates a highly lucrative market for a specialized data collection B2B SaaS platform focused specifically on the Indian subcontinent.¹⁹

Edge-Optimized Hardware Design for Low-Budget Telemetry

The core technological challenge in building a crowdsourced data collection network lies in designing a hardware device that is simultaneously affordable enough for mass consumer adoption and powerful enough to perform intelligent edge computing. The device must interface seamlessly with the vehicle's On-Board Diagnostics (OBD-II) port to capture internal telematics, utilize a Global Navigation Satellite System (GNSS) module for precise geospatial tagging, and employ high-definition cameras for visual perception.²¹

Microcontroller and Computational Hardware Analysis

Selecting the optimal processing unit requires a meticulous balancing of cost, computational capability, thermal management, and peripheral support. A comparative analysis of prevalent microcontrollers and single-board computers reveals distinct advantages and trade-offs for this specific data collection application.

Hardware Platform	Approximate Cost Bracket	AI Processing Capability	Video Encoding / Streaming Capability	Wireless Connectivity Integration	Suitability for Proposed Data Node
ESP32-S3 (AI CAM)	Very Low	Basic (Vector instructions suitable for lightweight neural networks)	Limited. Lacks H.264 hard encoding; relies on MJPEG. Struggles with continuous HD video storage. ²³	Wi-Fi / Bluetooth Built-in natively	Suboptimal. Best utilized for basic image classification or sending singular frames, rather than continuous HD video telemetry. ²⁴
Raspberry Pi Zero 2W	Low to Medium	Moderate (Quad-core ARM CPU architecture)	Good. Supports external camera modules	Wi-Fi / Bluetooth Built-in natively	Optimal. Strikes the perfect balance between

		e)	and efficient video logging to local mass storage. ²⁶		budget constraints, video handling, and peripheral integration via GPIO headers. ²⁸
NVIDIA Jetson Nano	High	Superior (128-core Maxwell GPU, specialized TensorRT support)	Excellent. Can run complex YOLO models and sophisticated deep learning pipelines locally. ³⁰	Requires external networking dongle	Overpowered and prohibitively expensive for a low-budget, mass-distributed consumer data collection node. ³²

The architectural evaluation demonstrates that the Raspberry Pi Zero 2W serves as the most viable and economically feasible core for the hardware prototype. While it lacks the dedicated tensor cores and immense parallel processing capabilities of the Jetson Nano, it possesses sufficient processing overhead to manage multiple sensor inputs simultaneously, encode video efficiently, and run lightweight edge-filtering algorithms without exceeding budget constraints.²⁶

Peripheral Integration and Telemetry Protocol

The hardware suite must be comprehensive in its data-gathering capabilities yet minimally invasive to the host vehicle. The visual component relies on a wide-angle, high-definition camera module capable of functioning in varying illumination scenarios, including intense glare and low-light conditions commonly encountered on Indian roads. Telematics data is extracted via a Controller Area Network (CAN) bus interface communicating directly with the vehicle's OBD-II port.²¹ This connection provides continuous access to granular mechanical and operational data, including engine RPM, precise vehicle speed, steering angle, and instantaneous braking pressure.²² Gathering this mechanical data is critical for correlating driver intent and vehicular response with the visual phenomena captured by the camera.³⁴

Geospatial intelligence is captured through a dedicated GNSS module, ensuring the accurate mapping of road anomalies, pothole locations, and dynamic traffic bottlenecks.²² To transmit this synthesized data to the cloud architecture, the system requires a robust cellular modem. A 4G LTE Hardware Attached on Top (HAT), such as the SIM7600E-H module, provides reliable broadband connectivity for uploading data.³⁵ Market research indicates that these specific 4G LTE HAT modules retail in the Indian market between INR 1900 and INR 5100, allowing the overall Bill of Materials (BOM) to remain within the constraints of a low-budget hardware project.³⁵

Artificial Intelligence Training and Active Learning Pipelines

The vast quantities of data aggregated by the hardware network form the foundation for a sophisticated B2B AI training pipeline. The primary objective is to refine foundation models to understand the distinct semantics and probabilistic nature of the Indian driving environment.

The Bandwidth Bottleneck and Active Learning

A critical infrastructural hurdle in cloud-based video analytics is the sheer volume of data generated by thousands of distributed camera nodes. The continuous transmission of high-definition video over 4G cellular networks is economically unviable due to exorbitant bandwidth costs and the inherent latency of mobile networks. Furthermore, transmitting hours of routine, uneventful highway driving yields minimal value for AI training, as standard models already perform adequately in these baseline scenarios.³⁹ The system must therefore implement intelligent, selective uploading mechanisms.

By utilizing the computational capabilities of the edge device, the platform can employ an Active Learning framework. Active learning algorithms dynamically select the most informative samples for annotation based on the AI model's internal uncertainty, thereby drastically reducing the volume of data required for effective training.⁴⁰ The Raspberry Pi Zero 2W can execute a lightweight, highly compressed model—such as a quantized YOLOv8 nano variant or a streamlined Contrastive Language-Image Pretraining (CLIP) model—solely to compute uncertainty scores or detect data drift in real-time.⁴³

When the edge model encounters a scenario it cannot confidently classify—such as a complex intersection deadlock, a severely degraded road surface, or an unrecognizable, hand-painted traffic sign—the system flags that specific temporal segment. The hardware then uploads only these high-value, high-uncertainty video clips to the cloud server.³⁹ This selective uploading approach minimizes bandwidth expenditure, optimizes cloud storage costs, and ensures that the central repository is populated exclusively with the rare, complex edge cases that are most valuable to autonomous driving developers.⁴⁸

Annotation Strategies for Indian-Specific Entities

The commercial value of the SaaS offering is directly proportional to the quality, accuracy, and specificity of its annotations. Datasets must be meticulously curated to include Indian-specific entities that are entirely absent from standard Western datasets like KITTI or nuScenes.² This includes diverse vehicle typologies such as cycle rickshaws, heavily loaded transport tractors, and customized local transport vehicles, alongside unpredictable pedestrian behaviors and non-standard infrastructure markers.⁵

Furthermore, the annotation process must capture dynamic human-vehicle interactions. In scenarios lacking functional traffic lights, intersection right-of-way is often negotiated via subtle eye contact and specific hand gestures between drivers, pedestrians, and traffic police. Advanced spatial tracking, semantic segmentation, and behavioral annotation must be applied to the collected video data to encode these culture-specific driving norms into the AI models.¹⁶ To facilitate VLM training, the annotations must evolve beyond simple bounding boxes to include comprehensive 3D scene graphs and Visual Question Answering (VQA) pairs.⁵² By generating rich textual descriptions of complex scenes, the platform enables VLMs to develop superior spatial understanding and reasoning capabilities, creating a highly specialized dataset that serves as the primary product licensed to Tier-1 automotive suppliers and OEMs.⁵⁴

Consumer Incentive Engineering: Driving Mass Adoption

The fundamental challenge of deploying a crowdsourced hardware network is overcoming the "cold start" problem: acquiring a critical mass of users to install the device and continuously share their driving data. Relying purely on altruism, technological enthusiasm, or vague promises of contributing to future technologies is demonstrably insufficient for achieving commercial scale. The platform must architect deeply compelling, tangible benefits that align directly with the daily economic, legal, and psychological drivers of Indian car commuters.

Decentralized Physical Infrastructure Networks (DePIN) and Tokenomics

The most potent mechanism for driving rapid network scaling among early adopters and gig-economy drivers is the application of blockchain-based tokenomics, specifically within the Decentralized Physical Infrastructure Network (DePIN) framework. Projects such as Hivemapper, DIMO, and NATIX have successfully demonstrated that tokenized reward systems can exponentially accelerate physical infrastructure deployment.⁵⁷ Hivemapper, utilizing proprietary dashcams and its native token, successfully mapped over 35% of global roads in a fraction of the time and cost it took centralized corporate entities like Google.⁵⁸

By integrating a similar tokenomic structure, daily commuters are transformed from passive

subjects into active, financially motivated stakeholders. Drivers install the hardware device and begin earning cryptographic tokens proportionally to the volume, quality, and novelty of the data they collect. The incentive algorithm can be dynamically weighted to reward driving in unexplored geographic regions, during challenging weather conditions, or in areas with known infrastructural deficits, thereby ensuring a diverse and comprehensive dataset.⁵⁹

The economic sustainability of this model relies on a burn-and-mint equilibrium. On the demand side, enterprise customers (automotive OEMs and AI researchers) purchase access to the driving data and AI models using fiat currency. These revenues are then utilized to buy back and burn the native tokens from the open market, creating deflationary pressure and stabilizing the token's value.⁶⁰ This creates a sustainable economic loop: as OEM demand for the data increases, the value of the rewards distributed to drivers increases, incentivizing further hardware adoption and denser data collection. This framework fundamentally reframes the act of driving from a sunk operational cost into a potential passive income stream.⁵⁹

Usage-Based Insurance (UBI) Integration

While tokenomics appeal heavily to early adopters and technologists, mainstream consumers generally require immediate, fiat-based financial benefits. Partnering with domestic Insurtech companies presents a massive opportunity to subsidize the hardware and reward safe driving with direct monetary savings. The regulatory environment in India is highly conducive to this model; in July 2022, the Insurance Regulatory and Development Authority of India (IRDAI) officially permitted the introduction of Usage-Based Insurance (UBI) products, including "Pay-As-You-Drive" (PAYD) and "Pay-How-You-Drive" (PHYD) models, as add-ons to standard motor policies.⁶²

Traditional insurance pricing relies on broad demographic generalizations and vehicle parameters, which inherently penalize safe drivers who subsidize the high-risk behavior of others. UBI, conversely, utilizes the granular telematics and visual data captured by the proposed hardware to assess individual risk profiles accurately.⁶² Insurtech startups such as Acko, Digit Insurance, and Navi, which operate on digital-first, asset-light models, are aggressively seeking innovative methods to reduce loss ratios and optimize claims processing.⁶⁴

By providing verified, tamper-proof telematics data to these forward-thinking insurers, the platform enables them to offer dynamic premium discounts to users. A driver who consistently maintains safe following distances, adheres to speed limits, and avoids harsh braking and rapid acceleration (as verified continuously by the edge device) can receive substantial reductions in their annual insurance premiums.⁶⁷ Furthermore, the granular data facilitates a shift from retrospective punishment to proactive coaching, utilizing gamification elements and driver scorecards to encourage safer behavior, which in turn reduces systemic accident rates.⁶¹ The prospect of saving thousands of rupees annually on insurance premiums serves as a highly

compelling, rational driver for hardware installation among mainstream commuters.

Dashcam Utility and Evidence Provision

Beyond sophisticated AI training and insurance integration, the device serves a fundamental, highly demanded utility for the Indian consumer: serving as an objective digital witness. With over 1.5 lakh road deaths reported annually in India, and disputes over traffic violations and insurance claims piling up in an overburdened legal system, demand for dashcams has shifted from a novelty to a necessity.⁷⁰ The Indian automotive camera market is experiencing rapid growth, projected to reach over USD 1 Billion by 2034, driven by a compound annual growth rate exceeding 12%.⁷¹

The proposed hardware directly serves this burgeoning market demand. In the event of an accident, the device acts as an objective, high-definition witness, expediting the insurance claims process and protecting the driver against fraudulent liability claims or unwarranted aggression from other drivers.⁷⁰ Furthermore, the presence of immutable video evidence combined with GNSS location data provides critical protection against false traffic challans (fines) issued by automated systems or traffic police. By fulfilling this basic, urgent need for legal and financial protection, the hardware device justifies its placement on the dashboard independently of its background data-collection functions.⁷⁰

On-Device Advanced Driver Assistance Systems (ADAS)

To ensure the hardware remains valuable even in rural areas with poor cellular connectivity, it must function as a standalone, active safety device. Utilizing the computational power of the Raspberry Pi Zero 2W, the system can deploy lightweight computer vision algorithms locally, providing immediate alerts to the driver without requiring a round-trip to cloud servers.

By employing open-source libraries such as OpenCV, the device's interior-facing camera (if configured as a dual-lens system) can monitor the driver's facial landmarks continuously. By calculating the Eye Aspect Ratio (EAR), tracking blink frequency, and analyzing mouth movements (yawning), the system can detect the onset of driver drowsiness or severe fatigue in real-time, sounding an immediate acoustic alarm.⁷⁶ Drowsiness is a leading factor in a significant percentage of road crashes, particularly among long-distance drivers and night-time commuters; mitigating this risk provides immense value.⁸⁰

Simultaneously, the forward-facing camera can run rudimentary object detection to identify severe, immediate road hazards. Advanced algorithms can detect massive, water-filled potholes, hidden speed breakers, or oncoming wrong-way traffic, providing the driver with critical advanced warnings that standard vehicle systems lack.⁷⁶ Offering these potentially life-saving features directly on the device creates immediate, non-monetary utility, drastically lowering the barrier to entry for safety-conscious consumers who may be unable to afford premium vehicles with factory-installed ADAS.⁷

Incentive Mechanism	Primary Driver Benefit	Mechanism of Value Creation	Target Demographic
DePIN Tokenomics	Passive Income Generation	Rewarded algorithmically for uploading high-value, rare edge-case video data. Demand driven by OEM SaaS subscriptions and token burn mechanisms.	Early adopters, gig-economy drivers, technology enthusiasts, fleet operators.
Usage-Based Insurance	Reduced Annual Premiums	Granular OBD-II telematics prove safe driving habits to Insurtech partners, lowering actuarial risk and justifying premium discounts.	Mainstream daily commuters, budget-conscious owners, young drivers seeking lower rates.
Video Evidence Provision	Legal Protection & Faster Payouts	Tamper-proof, timestamped video evidence eliminates ambiguity in accident liability assessments and protects against false traffic citations.	All vehicle owners navigating chaotic urban traffic.
On-Device ADAS Alerts	Immediate Collision Avoidance	Edge-processed OpenCV algorithms detect driver drowsiness (via facial landmarks) and structural road	Night-time commuters, long-distance travelers, drivers of older vehicles lacking modern safety

		hazards, sounding real-time acoustic alarms.	features.
--	--	--	-----------

Regulatory Compliance and Privacy-Preserving Architecture

In an era characterized by ubiquitous digital surveillance, modern consumers are increasingly hyper-aware of data privacy risks. Furthermore, any data collection enterprise operating within India must strictly adhere to the expansive requirements of the Digital Personal Data Protection (DPDP) Act of 2023. This comprehensive legislation mandates stringent consent architectures, strict purpose limitation, data minimization, and the implementation of reasonable security safeguards.⁸³ The rules operationalizing the Act mandate significant compliance obligations, which are being implemented in a phased manner between 2025 and 2027.⁸³

Consent Management and Privacy by Design

To build consumer trust and ensure absolute regulatory compliance, the hardware and software architecture must employ privacy-by-design principles at every level.⁸⁷ The system will incorporate a robust Consent Management System (CMS) allowing users granular, easily revocable control over what data is shared, with whom it is shared, and for what specific purposes, aligning directly with Section 6 of the DPDP Act.⁸⁸

More importantly, the edge-computing architecture inherently protects user privacy by ensuring that raw, unedited video does not indiscriminately flood central servers. The on-device processor will utilize localized machine learning models to dynamically blur faces, license plates, and other Personally Identifiable Information (PII) before the video packet is ever transmitted over the cellular network or stored on the device's SD card.⁹¹

For aggregate training tasks, the platform will explore the integration of Federated Learning (FL). FL is a distributed machine learning process where the edge devices download a global model, train it locally on the raw, sensitive driving data, and transmit only the updated mathematical parameters (weights and biases) back to the central server. The raw video and telematics data never leave the vehicle.⁹¹ By cryptographically and architecturally guaranteeing that the user's location history and personal identity cannot be reconstructed by the SaaS provider, state actors, or third parties, the platform mitigates privacy anxieties, transforming hesitant prospects into willing participants while remaining fully compliant with the DPDP Act.⁸⁷

B2B SaaS Feasibility and Market Strategy

The commercial engine sustaining this initiative is the B2B SaaS platform targeting the automotive sector. The Indian SaaS ecosystem is experiencing explosive growth, characterized by profound capital efficiency; research indicates that the vast majority of Indian SaaS companies operate with a burn multiple of less than 1.5x, proving that rapid scaling can be achieved with disciplined financial management.⁹⁵ Concurrently, the domestic automotive market is undergoing a rapid technological evolution, transitioning toward software-defined vehicles.⁹⁷

Major Indian OEMs, including Tata Motors, Mahindra & Mahindra, and Maruti Suzuki, are engaged in a highly competitive arms race to deliver semi-autonomous capabilities, connected-car features, and advanced safety suites to the mass market.⁹ However, as previously established, deploying generic Western ADAS stacks results in suboptimal performance and diminished consumer trust.⁷ These OEMs require continuous access to vast, localized datasets to retrain their perception algorithms and validate their safety models against the harsh realities of Indian infrastructure.¹¹

Tiered SaaS Product Offerings

The proposed business model monetizes the collected data through tiered B2B SaaS subscriptions, creating recurring, highly scalable revenue streams.

SaaS Subscription Tier	Product Offering	Target Customer Segment	Value Proposition
Tier 1: Raw Data Lake Access	Access to the vast repository of anonymized, unfiltered driving video and paired OBD-II telematics.	Large OEMs (Tata, Mahindra) and Tier-1 Suppliers (Bosch, ZF) with massive internal data science teams. ⁷²	Provides the raw volume needed for foundation model pre-training. Cheaper than deploying proprietary test fleets.
Tier 2: Annotated Edge-Case Datasets	Curated datasets featuring high-fidelity semantic segmentation, 3D scene graphs, and VQA pairs, categorized by	Autonomous driving startups (Minus Zero, Swaayatt Robots), AI research institutions. ⁹⁸	Drastically reduces time-to-market by providing ready-to-use, high-quality data for fine-tuning specific perception

	edge-case taxonomy (e.g., "monsoon weather anomalies").		modules.
Tier 3: Pre-Trained VLM APIs	Access to proprietary Vision-Language Models fully optimized for Indian traffic reasoning, available via API integration.	Mid-market automotive manufacturers, fleet management software providers. ³⁴	Allows companies without deep AI expertise to integrate state-of-the-art predictive reasoning into their software-defined vehicle architectures.

The recurring revenue generated from these long-term enterprise contracts provides the robust financial foundation required to sustain the commuter incentive programs, subsidize the hardware costs, and fund continuous algorithmic improvements, ensuring a highly defensible and self-reinforcing business cycle.

Hackathon Presentation Strategy (Vihaan 9.0 IEEE DTU)

Success at the Vihaan 9.0 hackathon necessitates a highly strategic alignment with the event's structure, time constraints, and specific judging rubrics. The 24-hour hardware track is an intensive crucible that demands flawless execution and meticulous time management across three distinct 8-hour phases: Ideation, Building, and Testing/Iteration.¹⁰⁰ Furthermore, the judging criteria heavily weight technical innovation (15%), market feasibility (15%), financial feasibility (15%), and the clarity and professionalism of the final presentation.¹⁰¹

Execution Timeline: The 24-Hour Roadmap

To maximize efficiency, the team must adhere to a strict operational timeline that balances hardware prototyping with the development of a compelling business narrative.

Phase 1: Ideation and Architectural Finalization (Hours 1-8)

The initial phase must be dedicated to finalizing the hardware schematics and establishing the software repository. Given the extreme time constraints, the team will avoid writing low-level drivers from scratch, instead utilizing pre-existing, open-source computer vision libraries (such as OpenCV or pre-trained YOLOv8 nano models) to expedite development.⁴⁵ The architectural blueprint must clearly delineate the communication protocols between the Raspberry Pi Zero

2W, the OBD-II CAN bus reader, and the camera module. During this phase, the team will also draft the core business logic for the active learning filter, ensuring the device is programmed to only save frames where object detection confidence thresholds drop below a predetermined limit.

Phase 2: Hardware Assembly and Core Programming (Hours 9-16)

This is the intensive execution phase. The physical prototype must be rapidly assembled, ensuring robust electrical connections between the GPIO pins, the camera ribbon cable, and the 4G LTE HAT. The primary programming objective is to establish the fundamental data pipeline: successfully capturing a video frame, reading simultaneous telematics data from a simulated OBD-II stream, running the lightweight object detection model, and saving the correlated, timestamped data packet to local storage. Advanced cloud connectivity features and complex blockchain integrations should be deprioritized or mocked up until the local edge processing pipeline is demonstrably functional.

Phase 3: Testing, Iteration, and Pitch Refinement (Hours 17-24)

The final phase focuses on debugging the hardware and crafting the presentation narrative. The prototype must be rigorously tested under simulated conditions. To demonstrate functionality within a controlled environment, the team will point the hardware camera at a high-resolution monitor displaying footage of chaotic Indian traffic, proving that the active learning trigger functions correctly when presented with visual anomalies. Simultaneously, the presentation deck must be developed to directly address every aspect of the judging criteria, ensuring that technical prowess is matched by business acumen.¹⁰²

Pitch Narrative Structure and Delivery

The pitch must not merely demonstrate a working gadget; it must sell a comprehensive, highly lucrative business ecosystem. The presentation should be structured to immediately capture attention and sequentially build a compelling case for investment and adoption.

1. **The Hook (Addressing Innovation & Problem Identification):** Open the presentation with a stark, visual demonstration of standard ADAS failing on Indian roads due to chaotic edge cases (e.g., a sudden influx of cattle or faded lane markings). Establish the foundational premise that the barrier to autonomy in emerging markets is a critical lack of localized data, not algorithmic capability.³
2. **The Technical Solution (Addressing Technical Feasibility):** Reveal and demonstrate the physical hardware prototype. Explain the elegance and cost-efficiency of using low-budget edge hardware (Raspberry Pi) combined with sophisticated active learning algorithms to solve the massive bandwidth and storage bottlenecks inherent in continuous video transmission.⁴¹ Highlight the critical integration of OBD-II telematics to provide deep context to the visual data.
3. **The Commercial Engine (Addressing Financial & Market Feasibility):** Outline the dual-sided nature of the market. Explain the highly scalable B2B SaaS model targeting

desperate automotive OEMs willing to pay a premium for localized VLM training data.⁹⁷

Immediately counter expected skepticism regarding user adoption by detailing the robust trifecta of consumer incentives: DePIN tokenomics for passive income, UBI integration for massive insurance discounts, and the provision of free on-device ADAS features.⁶⁰

4. **Regulatory Awareness (Addressing Adaptability and Risk Management):** Briefly highlight the system's privacy-by-design architecture, emphasizing automated edge-blurring and potential federated learning applications to ensure strict compliance with India's DPDP Act 2023.⁸⁴ This demonstrates to the judges a mature, enterprise-level understanding of real-world deployment challenges and legal liabilities.
5. **The Vision (Addressing Impact & Potential):** Conclude the pitch by framing the project not just as a profitable data collection tool, but as the foundational digital infrastructure required to democratize autonomous safety features and drastically reduce traffic fatalities for the next billion drivers in emerging markets worldwide.⁹⁸

Strategic Conclusions

The pursuit of autonomous driving in the Indian subcontinent requires a radical departure from the highly structured, deterministic methodologies currently employed in Western markets. The chaotic, probabilistic nature of Indian roads demands the advanced contextual reasoning capabilities inherent in modern Vision-Language Models. However, these advanced models exhibit a ravenous appetite for highly specific, localized training data featuring complex edge cases.

The proposed platform elegantly bridges this critical data gap by deploying a decentralized network of low-cost, intelligent hardware nodes directly into the vehicles of daily commuters. By processing data at the edge and employing sophisticated active learning algorithms, the system minimizes exorbitant transmission and storage costs while accurately capturing the most elusive and valuable traffic anomalies.

Crucially, the success of this hardware network is entirely dependent on consumer adoption, which is secured through a robust, multi-faceted incentive structure. By offering tangible financial rewards via tokenized gamification and Usage-Based Insurance discounts, alongside immediate on-device safety utilities and strict adherence to the privacy mandates of the DPDP Act 2023, the platform transforms ordinary, potentially skeptical commuters into willing, active partners in the AI revolution. Ultimately, this comprehensive approach not only constructs a highly lucrative, scalable B2B SaaS ecosystem for the automotive sector but also lays the critical data infrastructure necessary to realize safe, intelligent, and autonomous mobility in the world's most complex and challenging driving environments.

Works cited

1. "Startups [Swaayatt, Minus Zero, RoshAI] Say India Is Ideal for Testing Self-Driving Cars" : r/reinforcementlearning - Reddit, accessed March 26, 2026,

- https://www.reddit.com/r/reinforcementlearning/comments/1cg4cu9/startups_sw_aayatt_minus_zero_roshai_say_india_is/
2. DriveIndia: An Object Detection Dataset for Diverse Indian Traffic Scenes - arXiv, accessed March 26, 2026, <https://arxiv.org/html/2507.19912v1>
 3. Why ADAS Is Being Disabled in India And How Starkenn Is Solving This Real Problem, accessed March 26, 2026, <https://www.starkenn.com/post/why-adas-is-being-disabled-in-india-and-how-s-tarkenn-is-solving-this-real-problem>
 4. Why does ADAS Stand Unsuccessful on Indian Roads? - GoMechanic, accessed March 26, 2026, <https://gomechanic.in/blog/why-does-adas-stand-unsuccessful-on-indian-roads/>
 5. Indian Driving Dataset-Detections (YOLOv11) - Kaggle, accessed March 26, 2026, <https://www.kaggle.com/datasets/redzapdos123/indian-driving-dataset-detections-yolov11>
 6. Real-World ADAS Comparison Test | Maruti vs Tata vs Mahindra vs Kia vs Hyundai vs Toyota | autoX, accessed March 26, 2026, <https://www.youtube.com/watch?v=C4X0fTwTx5Y>
 7. ADAS Explained: Features Involved And The Challenges It Will Face In India - CarDekho, accessed March 26, 2026, <https://www.cardekho.com/india-car-news/what-is-adas-how-does-it-help-and-what-are-the-challenges-it-faces-in-india-28054.htm>
 8. Mahindra XUV700 ADAS technology explained - CarWale, accessed March 26, 2026, <https://www.carwale.com/news/mahindra-xuv700-adas-technology-explained/>
 9. Adas Safety Features: When the weather turns hostile, how smart are 'smart' vehicles?, ETAuto, accessed March 26, 2026, <https://auto.economictimes.indiatimes.com/news/auto-technology/when-the-weather-turns-hostile-how-smart-are-smart-vehicles/128147637>
 10. Market Research on Actual User Experience and Efficacy of ADAS Features in Indian Driving Conditions 2026-26-0035 - SAE International, accessed March 26, 2026, <https://www.sae.org/papers/market-research-actual-user-experience-efficacy-adas-features-indian-driving-conditions-2026-26-0035>
 11. India's Road to Autonomous Driving - From Chaos to Opportunity - EV Tech News, accessed March 26, 2026, <https://evtechnews.in/indias-road-to-autonomous-driving-from-chaos-to-opportunity/>
 12. [2412.14446] VLM-AD: End-to-End Autonomous Driving through Vision-Language Model Supervision - arXiv, accessed March 26, 2026, <https://arxiv.org/abs/2412.14446>
 13. AutoDriDM: An Explainable Benchmark for Decision-Making of Vision-Language Models in Autonomous Driving - arXiv, accessed March 26, 2026, <https://arxiv.org/html/2601.14702v1>
 14. (PDF) Explainable Artificial Intelligence for Autonomous Driving: A Comprehensive Overview and Field Guide for Future Research Directions - ResearchGate,

accessed March 26, 2026,

https://www.researchgate.net/publication/382403041_Explainable_Artificial_Intelligence_for_Autonomous_Driving_A_Comprehensive_Overview_and_Field_Guide_for_Future_Research_Directions

15. DriveLLaVA: Human-Level Behavior Decisions via Vision Language Model - PMC, accessed March 26, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC11243790/>
16. Authorized Traffic Controller Hand Gesture Recognition for Situation-Aware Autonomous Driving - MDPI, accessed March 26, 2026, <https://www.mdpi.com/1424-8220/21/23/7914>
17. CoT-VLM4Tar: Chain-of-Thought Guided Vision-Language Models for Traffic Anomaly Resolution - arXiv, accessed March 26, 2026, <https://arxiv.org/html/2503.01632v1>
18. VERDI: VLM-Embedded Reasoning for Autonomous Driving - arXiv, accessed March 26, 2026, <https://arxiv.org/html/2505.15925v1>
19. IDD-3D: Indian Driving Dataset for 3D Unstructured Road Scenes, accessed March 26, 2026, <https://idd3d.github.io/>
20. India Driving Dataset, accessed March 26, 2026, <https://idd.insaan.iiit.ac.in/>
21. ESP32 Vehicle Telematics - OBD2 (CAN) and GPS Reporting System - Instructables, accessed March 26, 2026, <https://www.instructables.com/ESP32-Vehicle-Telematics-OB2-CAN-and-GPS-Reportin/>
22. OBD-II Data Logger with IoT Dashboard and Dashcam - Circuit Digest, accessed March 26, 2026, <https://circuitdigest.com/microcontroller-projects/obd-ii-data-logger-with-iot-dashboard-and-dashcam>
23. [ESP32-S3 AI Smart Camera Module] Implementing Video Recording: MJPEG Encoding and VLC Analysis - Edge AI - EEWorld, accessed March 26, 2026, <https://en.eeworld.com.cn/bbs/thread-1316911-1-1.html>
24. Capture Video on Vision AI V2 and Run Object Detection Simultaneously - Grove, accessed March 26, 2026, <https://forum.seeedstudio.com/t/capture-video-on-vision-ai-v2-and-run-object-detection-simultaneously/294827>
25. Building Smarter Camera Applications with esp-video - Espressif Developer Portal, accessed March 26, 2026, <https://developer.espressif.com/blog/2025/09/esp-video-introduction/>
26. Building a Simple Raspberry Pi Dashcam, accessed March 26, 2026, <https://era86.github.io/2018/04/10/pi-zero-w-usb-dashcam.html>
27. How to Make a Dashcam with a Raspberry Pi - YouTube, accessed March 26, 2026, <https://www.youtube.com/watch?v=aztYb1XLp2Y>
28. 24 Best Raspberry Pi Camera Projects to Help You Get Started! - Seeed Studio, accessed March 26, 2026, <https://www.seeedstudio.com/blog/2020/01/20/20-best-raspberry-pi-camera-projects-to-try-in-2020/>
29. building a dashcam - Raspberry Pi Forums, accessed March 26, 2026, <https://forums.raspberrypi.com/viewtopic.php?t=331222>

30. Jetson Nano vs Raspberry Pi AI: The Ultimate Performance Comparison fo - ThinkRobotics.com, accessed March 26, 2026,
<https://thinkrobotics.com/blogs/learn/jetson-nano-vs-raspberry-pi-ai-the-ultimate-performance-comparison-for-edge-computing>
31. Jetson Nano vs Raspberry Pi: Which Is Better for AI Projects - twowin technology, accessed March 26, 2026,
<https://twowintech.com/jetson-nano-vs-raspberry-pi-which-is-better-for-ai-projects/>
32. Introduction and difference analysis of Raspberry Pi and Jetson Nano self-driving robot cars, accessed March 26, 2026,
<https://www.xiaorgeek.net/blogs/news/introduction-and-difference-analysis-of-raspberry-pi-and-jetson-nano-self-driving-robot-cars>
33. Build a Dash Cam with a Raspberry Pi, Two Cameras, and GPS | Lifehacker, accessed March 26, 2026,
<https://lifehacker.com/build-a-dash-cam-with-a-raspberry-pi-two-cameras-and-1730114347>
34. Implementing Vehicle Telematics with Raspberry Pi and PiCAN FD HAT - Copperhill, accessed March 26, 2026,
<https://copperhilltech.com/blog/implementing-vehicle-telematics-with-raspberry-pi-and-pican-fd-hat/>
35. SIM7600E-H 4G HAT for Raspberry Pi - IndiaMART, accessed March 26, 2026,
<https://www.indiamart.com/proddetail/sim7600e-h-4g-hat-for-raspberry-pi-2851976485688.html>
36. Buy Waveshare SIM7600G-H 4G HAT For Raspberry Pi, LTE Cat-4 4G / 3G / 2G Support online - Robu.in, accessed March 26, 2026,
<https://robu.in/product/waveshare-sim7600g-h-4g-hat-for-raspberry-pi-lte-cat-4-4g-3g-2g-support/>
37. GeeekPi RPi 4G 3G Hat for Raspberry Pi ZeroZero India | Ubuy, accessed March 26, 2026,
<https://www.ubuy.co.in/product/8AQWSKPM-geeekpi-rpi-4g-3g-hat-for-raspberry-pi-zero-zero-w-zero-wh-2b-3b-3b-4b-jetson-nano-supports-sms-mms->
38. Raspberry Pi 4G / 3G Hat For Pi Zero/W/WH/2B/3B+/4B,Jetson Nano Trader - Retailer from New Delhi, accessed March 26, 2026,
<https://www.call4electronics.in/raspberry-pi.html>
39. Bandwidth Allocation for Cloud-Augmented Autonomous Driving - arXiv, accessed March 26, 2026, <https://arxiv.org/html/2503.20127v1>
40. How Can Active Learning Help Train Autonomous Vehicles? - NVIDIA Blog, accessed March 26, 2026, <https://blogs.nvidia.com/blog/what-is-active-learning/>
41. Advanced Active Learning Strategies for Object Detection - mediaTUM, accessed March 26, 2026,
https://mediatum.ub.tum.de/doc/1585225/d0g2bf2vp7u14v28cj7elg618.Advanced_Active_Learning_Strategies_for_Object_Detection.pdf
42. Lecture 26: Active Learning for Network Training: Uncertainty Sampling and other approaches. - YouTube, accessed March 26, 2026,
<https://www.youtube.com/watch?v=fwHZtqr-uBY>

43. Vision-Language Models for Autonomous Driving: CLIP-Based Dynamic Scene Understanding - MDPI, accessed March 26, 2026, <https://www.mdpi.com/2079-9292/14/7/1282>
44. Edge-Cloud Collaborative Motion Planning for Autonomous Driving with Large Language Models - arXiv, accessed March 26, 2026, <https://arxiv.org/html/2408.09972v1>
45. Lane and Traffic Sign Detection for Autonomous Vehicles: Addressing Challenges on Indian Road Conditions - PMC, accessed March 26, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC11795139/>
46. Object Detection and Classification Framework for Analysis of Video Data Acquired from Indian Roads - MDPI, accessed March 26, 2026, <https://www.mdpi.com/1424-8220/24/19/6319>
47. A Survey on Video Anomaly Detection via Deep Learning: Human, Vehicle, and Environment - arXiv, accessed March 26, 2026, <https://arxiv.org/html/2508.14203v1>
48. Active Learning from Scene Embeddings for End-to-End Autonomous Driving - arXiv, accessed March 26, 2026, <https://arxiv.org/html/2503.11062v1>
49. Active Learning, Data Selection, Data Auto-Labeling, and Simulation in Autonomous Driving — Part 3 - Isaac Kargar, accessed March 26, 2026, <https://kargarisaac.medium.com/active-learning-data-selection-data-auto-labeling-and-simulation-in-autonomous-driving-part-3-24e2711e2ba3>
50. DATS_2022: A versatile indian dataset for object detection in unstructured traffic conditions, accessed March 26, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC9309657/>
51. Data Annotation Types to execute Autonomous Driving - INDIAai, accessed March 26, 2026, <https://indiaai.gov.in/article/data-annotation-types-to-execute-autonomous-driving>
52. NuScenes-SpatialQA: A Spatial Understanding and Reasoning Benchmark for Vision-Language Models in Autonomous Driving - CVF Open Access, accessed March 26, 2026, https://openaccess.thecvf.com/content/ICCV2025W/WDFM-AD/papers/Tian_NuScenes-SpatialQA_A_Spatial_Understanding_and_Reasoning_Benchmark_for_Vision-Language_Models_ICCVW_2025_paper.pdf
53. Indian Traffic VQA Dataset - Kaggle, accessed March 26, 2026, <https://www.kaggle.com/datasets/chandramohanbhuma/indian-traffic-vqa-dataset>
54. Semantic VLM Dataset for Safe Autonomous Driving - arXiv, accessed March 26, 2026, <https://arxiv.org/html/2511.10701v1>
55. AniDriveQA: a VQA dataset for driving scenes with animal presence - Frontiers, accessed March 26, 2026, <https://www.frontiersin.org/journals/robotics-and-ai/articles/10.3389/frobt.2025.1684845/full>
56. Drive-Bench, accessed March 26, 2026, <https://drive-bench.github.io/>
57. Projects like Hivemapper and DIMO to Replace Google Maps and Uber by 2030?

- Waivio, accessed March 26, 2026,
<https://www.waivio.com/@brennanhm/projects-like-hivemapper-and-dimo-to-replace-google-maps-and-uber-by-2030>
- 58. Here's why Solana-based Hivemapper thinks it can beat Google Maps - DL News, accessed March 26, 2026,
<https://www.dlnews.com/articles/defi/solana-depin-project-hivemapper-aims-to-beat-google-maps/>
- 59. Top 5 Drive-to-Earn Projects You Can Use in Your Car Today! - YouTube, accessed March 26, 2026, <https://www.youtube.com/watch?v=ICEfj7byQ5k>
- 60. Case Study: Reimagining Maps through Blockchain Incentives — The Hivemapper Story | by Hilary H Brown | Medium, accessed March 26, 2026,
<https://medium.com/@hilary.h.brown/case-study-reimagining-maps-through-blockchain-incentives-the-hivemapper-story-9c31753287d1>
- 61. Driver incentive programs, gamification and rewards | Verizon Connect, accessed March 26, 2026,
<https://www.verizonconnect.com/resources/article/driver-incentive-program-gamification/>
- 62. Usage-Based Car Insurance in India: A Comprehensive Guide - Policybazaar.com, accessed March 26, 2026,
<https://www.policybazaar.com/motor-insurance/car-insurance/articles/usage-based-car-insurance-in-india/>
- 63. Usage-Based Insurance: The case for Motor Insurance | by Sahil Dhok - Medium, accessed March 26, 2026,
<https://medium.com/insurance-2030/usage-based-insurance-the-case-for-motor-insurance-1701bb345ac6>
- 64. Navi forays into motor insurance, set to compete with Acko and Digit - Medial App, accessed March 26, 2026,
<https://medial.app/news/navi-forays-into-motor-insurance-set-to-compete-with-acko-and-digit-85dc34079329a>
- 65. Navi forays into motor insurance, set to compete with Acko and Digit - Entrackr, accessed March 26, 2026,
<https://entrackr.com/snippets/navi-forays-into-motor-insurance-set-to-compete-with-acko-and-digit-11221163>
- 66. Insurtech Startups to Watch in India - BillCut, accessed March 26, 2026,
<https://www.billcut.com/blogs/insurtech-startups-india/>
- 67. Are You a Good Driver? Sharing Data on Your Driving Could Lower Your Car Insurance Rates | Money, accessed March 26, 2026,
<https://money.com/telematics-savings-on-car-insurance/>
- 68. How Does Usage-Based Insurance Work?, accessed March 26, 2026,
<https://www.indusindinsurance.com/insurance/knowledge-center/insurance-reads/how-does-usage-based-car-insurance-work.aspx>
- 69. Usage Based Insurance (UBI) Global News - PTOLEMUS Consulting Group, accessed March 26, 2026,
<https://www.ptolemus.com/usage-based-insurance-global-news/>
- 70. From novelty to necessity: How dash cams in India are becoming a standard car

accessory, accessed March 26, 2026,

<https://m.economictimes.com/industry/auto/auto-news/from-novelty-to-necessity-how-dash-cams-in-india-are-becoming-a-standard-car-accessory/articleshow/123829617.cms>

71. India Automotive Camera Market Size, Statistics Report 2034 - IMARC, accessed March 26, 2026, <https://www.imarcgroup.com/india-automotive-camera-market>
72. India Automotive Camera Market Size, Share | Report, 2032 - MarketsandMarkets, accessed March 26, 2026, <https://www.marketsandmarkets.com/Market-Reports/india-automotive-camera-market-221377397.html>
73. India Automotive Camera Market - Size, Growth & Analysis - Mordor Intelligence, accessed March 26, 2026, <https://www.mordorintelligence.com/industry-reports/india-automotive-camera-market>
74. How to Use Dashcam Footage for Faster Insurance Claims?, accessed March 26, 2026, <https://www.acko.com/car-insurance/use-of-dashcam-for-faster-claims/>
75. India Dashcams Market Size, Growth, Share and Forecast 2032 - Credence Research, accessed March 26, 2026, <https://www.credenceresearch.com/report/india-dashcams-market>
76. Monitoring System of Drowsiness and Lost Focused Driver Using Raspberry Pi - PMC - NIH, accessed March 26, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC7898092/>
77. Enhancing Road Safety: Deep Learning-Based Intelligent Driver Drowsiness Detection for Advanced Driver-Assistance Systems - MDPI, accessed March 26, 2026, <https://www.mdpi.com/2079-9292/13/4/708>
78. Driver Drowsiness Detector System using Raspberry Pi and OpenCV - Circuit Digest, accessed March 26, 2026, <https://circuitdigest.com/microcontroller-projects/driver-drowsiness-detector-using-raspberry-pi-and-opencv>
79. Raspberry Pi-Based Driver Tiredness Monitoring and Alerting System for Truck, accessed March 26, 2026, <https://rsisinternational.org/journals/ijriss/articles/raspberry-pi-based-driver-tiredness-monitoring-and-alerting-system-for-truck/>
80. Raspberry Pi Meets Road Safety: Launch of Smart Drowsiness Detection System - Issuewire, accessed March 26, 2026, <https://www.issuewire.com/raspberry-pi-meets-road-safety-launch-of-smart-drowsiness-detection-system-1835180557760830>
81. Potholes and traffic signs detection by classifier with vision transformers - PMC - NIH, accessed March 26, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC10817940/>
82. iWatchRoad: Scalable Detection and Geospatial Visualization of Potholes for Smart Cities, accessed March 26, 2026, <https://arxiv.org/html/2508.10945v1>
83. India's Digital Personal Data Protection Act 2023 brought into force - Hogan Lovells, accessed March 26, 2026, <https://www.hoganlovells.com/en/publications/indias-digital-personal-data-prote>

[ction-act-2023-brought-into-force-](#)

84. DPDP Rules 2025- Analysis of Industry implications and Compliance Guidance, accessed March 26, 2026, <https://www.dpdpa.com/blogs/DPDP%20Rules%202025-%20Analysis%20of%20Industry%20implications.html>
85. The Impact of India's New Digital Personal Data Protection Rules - Privacy World, accessed March 26, 2026, <https://www.privacyworld.blog/2025/04/the-impact-of-indias-new-digital-personal-data-protection-rules/>
86. DPDPA Rules 2025: Complete Compliance Guide for Indian Startups and Digital Businesses, accessed March 26, 2026, <https://thestartupzone.in/legal-documentation/dpdpa-rules-2025-complete-compliance-guide-for-indian-startups-and-digital-businesses/>
87. Integrating Cybersecurity and Data Protection under the DPDP Act, 2023 - Informatics, accessed March 26, 2026, <https://informatics.nic.in/files/websites/october-2025/cybersecurity-and-privacy-in-governance.php>
88. India DPDP Phase 2 Compliance Guide (Digital Personal Data Protection Act), accessed March 26, 2026, <https://secureprivacy.ai/blog/india-dpdp-phase-2>
89. India publishes consent management rules under Digital Personal Data Protection Act, accessed March 26, 2026, <https://www.hoganlovells.com/en/publications/india-publishes-consent-management-rules-under-digital-personal-data-protection-act>
90. Decoding India's draft DPDPA rules for the world - IAPP, accessed March 26, 2026, <https://iapp.org/news/a/decoding-india-s-draft-dpdpa-rules-for-the-world>
91. A Comprehensive Evaluation of Privacy-Preserving Mechanisms in Cloud-Based Big Data Analytics: Challenges and Future Research Directions - Preprints.org, accessed March 26, 2026, <https://www.preprints.org/manuscript/202601.1025>
92. Privacy-Preserving AI at the Edge - XenonStack, accessed March 26, 2026, <https://www.xenonstack.com/blog/privacy-preserving-ai-edge>
93. Privacy Issues in Edge Computing, accessed March 26, 2026, <https://www.cs.wm.edu/~liqun/paper/book-privacy-21.pdf>
94. Analysis of Privacy-Preserving Edge Computing and Internet of Things Models in Healthcare Domain - PMC, accessed March 26, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC8739545/>
95. B2B SaaS marketing: Expand beyond India - Think with Google APAC, accessed March 26, 2026, <https://business.google.com/in/think/future-of-marketing/saas-business-growth-strategy-india/>
96. Bellwethers of Indian SaaS | EY, accessed March 26, 2026, <https://www.ey.com/content/dam/ey-unified-site/ey-com/en-in/insights/technology/documents/ey-bellwethers-of-indian-saas.pdf>
97. India Connected Car Market - Companies, Analysis, Share & Trends - Mordor Intelligence, accessed March 26, 2026, <https://www.mordorintelligence.com/industry-reports/india-connected-car-mark>

[et](#)

98. Minus Zero: Full Stack AI for ADAS/AD, accessed March 26, 2026,
<https://minuszero.ai/>
99. India's Road Challenges Spark Innovations In Autonomous Driving Globally - DriveSpark, accessed March 26, 2026,
<https://www.drivespark.com/four-wheelers/2024/indias-road-challenges-autonomous-driving-innovation-011-052355.html>
100. Vihaan 9.0 : Hardware Track - 2026 - Unstop, accessed March 26, 2026,
<https://unstop.com/hackathons/vihaan-90-hardware-track-vihaan-90-ieee-dtu-dtu-new-delhi-1664370>
101. #IETHackathon – Dirty Data Judging Criteria The Judging Panel will be evaluating each entry based on the technical feasibility, accessed March 26, 2026,
<https://www.theiet.org/media/10197/judging-criteria-young-professionals-hackathon.pdf>
102. Judging Criteria and Scoring - Innovation Competition - TIC Americas, accessed March 26, 2026,
<https://ticamericas.net/cic/rules/judging-criteria-and-scoring/>
103. Hackathon judging: 6 criteria to pick winning projects - TAIKAI Network, accessed March 26, 2026, <https://taikai.network/en/blog/hackathon-judging>
104. Understanding hackathon submission and judging criteria: What really counts - Devpost, accessed March 26, 2026,
<https://info.devpost.com/blog/understanding-hackathon-submission-and-judging-criteria>